

Outline

Overview of the SQL Query Language
Data Definition
Basic Query Structure
Additional Basic Operations
Set Operations
Null Values
Aggregate Functions
Nested Subqueries
Modification of the Database

Chapter 3: Introduction to SQL

History

IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory

(Sequel: A structured English query language, 1974)

Renamed SQL (Structured Query Language)

ANSI and ISO standard SQL:

SQL-86, SQL-89, SQL-92

SQL:1999, SQL:2003, SQL:2006, SQL:2008, SQL:2011, SQL:2016, SQL:2019

The benefits of the standards:

Points out which language extensions are important and useful.

Guides the development of SQL implementations in databases.

Provides the common syntax and semantics that most databases will implement.

Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.

Data Definition Language

The SQL data-definition language (DDL) allows the specification of information about relations, including:

The schema for each relation.

The domain of values associated with each attribute.

Integrity constraints

And as we will see later, also other information such as

The set of indices to be maintained for each relations.

Security and authorization information for each relation.

The physical storage structure of each relation on disk.

Domain Types in SQL

char(n). Fixed length character string, with user-specified length n .

varchar(n). Variable length character strings, with user-specified maximum length n .

int. Integer (a finite subset of the integers that is machine-dependent).

smallint. Small integer (a machine-dependent subset of the integer domain type).

numeric(p,d). Fixed point number, with user-specified precision of p digits, with d digits to the right of decimal point.

→ number(3,1) allows 44.5 to be stored exactly, but neither 444.5 or 0.32

real, double precision. Floating point and double-precision floating point numbers, with machine-dependent precision.

float(n). Floating point number, with user-specified precision of at least n digits.

Built-in Data Types in SQL

date: Dates, containing a (4 digit) year, month and date

Example: `date '2005-7-27'`

time: Time of day, in hours, minutes and seconds.

Example: `time '09:00:30'` `time '09:00:30.75'`

timestamp: date plus time of day

Example: `timestamp '2005-7-27 09:00:30.75'`

interval: period of time

Example: `interval '1' day`

Subtracting a date/time/timestamp value from another gives an interval value

Interval values can be added to date/time/timestamp values

date, time functions:

`current_date()`, `current_time()`

`year(x)`, `month(x)`, `day(x)`, `hour(x)`, `minute(x)`, `second(x)`

Create Table Construct

An SQL relation is defined using the `create table` command:

```
create table r (A1 D1, A2 D2, ..., An Dn,
               (integrity-constraint1),
               ...,
               (integrity-constraintk))
```

r is the name of the relation

each A_i is an attribute name in the schema of relation r

D_i is the data type of values in the domain of attribute A_i

Example:

```
create table instructor (
    ID      char(5),
    name    varchar(20) not null,
    dept_name varchar(20),
    salary   numeric(8,2))
```

`insert into instructor values ('10211', 'Smith', 'Biology', 66000);`

`insert into instructor values ('10211', null, 'Biology', 66000);`

not null

primary key (A_1, \dots, A_n)

foreign key (A_m, \dots, A_n) **references** r

Example: Declare ID as the primary key for `instructor`

```
create table instructor (
    ID      char(5),
    name    varchar(20) not null,
    dept_name varchar(20),
    salary   numeric(8,2),
    primary key (ID),
    foreign key (dept_name) references department)
```

primary key declaration on an attribute automatically ensures **not null**

And a Few More Relation Definitions

```

create table student (
    ID      varchar(5),
    name    varchar(20) not null,
    dept_name varchar(20),
    tot_cred numeric(3,0) default 0,
    primary key (ID),
    foreign key (dept_name) references department );

```

create table takes (

```

    ID      varchar(5),
    course_id  varchar(8),
    sec_id    varchar(8),
    semester  varchar(6),
    year     numeric(4,0),
    grade    varchar(2),
    primary key (ID, course_id, sec_id, semester, year),
    foreign key (ID) references student,
    foreign key (course_id, sec_id, semester, year) references section );

```

Note: **sec_id** can be dropped from primary key above, to ensure a student cannot be registered for two sections of the same course in the same semester

And more still

```

create table course (
    course_id   varchar(8) primary key,
    title        varchar(50),
    dept_name   varchar(20),
    credits      numeric(2,0),
    foreign key (dept_name) references department (dept_name);

foreign key (dept_name) references department
on delete cascade |set null |restrict |set default
on update cascade |set null |restrict |set default,

```

teaches

ID	course_id
10101	CS-101
12121	FIN-201
76766	BIO-101

instructor

ID	name	dept_name
10101	Srinivasan	Comp. Sci.
12121	Wu	Finance
76766	Mozart	Music

Drop and Alter Table Constructs

drop table student
 Deletes the table and its contents

delete from student
 Deletes all contents of table, but retains table

alter table

- alter table r add A D**
 - where A is the name of the attribute to be added to relation r and D is the domain of A.
 - All tuples in the relation are assigned *null* as the value for the new attribute.
 - alter table student add resume varchar(256);**
- alter table r drop A**
 - where A is the name of an attribute of relation r
 - Dropping of attributes not supported by many databases

*SQL and Relational Algebra

```

select A1, A2, ... An
from r1, r2, ..., rm
where P

```

is equivalent to the following expression in multiset relational algebra

$$\prod_{A_1, \dots, A_n} (\sigma_P (r_1 \times r_2 \times \dots \times r_m))$$

*SQL and Relational Algebra

Basic Query Structure

The SQL **data-manipulation language (DML)** provides the ability to query information, and insert, delete and update tuples

A typical SQL query has the form:

```

select A1, A2, ..., An
from r1, r2, ..., rm
where P

```

A_i represents an attribute

R_i represents a relation

P is a predicate.

The result of an SQL query is a relation.

The select Clause

The **select** clause lists the attributes desired in the result of a query
 corresponds to the **projection operation** of the relational algebra

Example: find the names of all instructors:

```

select name
from instructor

```

NOTE: SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)

E.g. $Name \equiv NAME \equiv name$

Some people use upper case wherever we use bold font.

The select Clause (Cont.)

SQL allows duplicates in relations as well as in query results.
 To force the elimination of duplicates, insert the keyword **distinct** after select.

Find the names of all departments with instructor, and remove duplicates

```

select distinct dept_name
from instructor

```

The keyword **all** specifies that duplicates not be removed.

```

select all dept_name
from instructor

```

The select Clause (Cont.)

An asterisk in the select clause denotes "all attributes"

```
select *
from instructor
```

The **select** clause can contain arithmetic expressions involving the operation, +, -, *, and /, and operating on constants or attributes of tuples.

The query:

```
select ID, name, salary/12
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

The where Clause

The **where** clause specifies conditions that the result must satisfy

Corresponds to the **selection predicate** of the relational algebra.

To find all instructors in Comp. Sci. dept with salary > 80000

```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary > 80000
```

Comparison results can be combined using the logical connectives **and**, **or**, and **not**.

Comparisons can be applied to results of arithmetic expressions.

Where Clause Predicates

SQL includes a **between** comparison operator

Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, $\geq \$90,000$ and $\leq \$100,000$)

```
select name
from instructor
where salary between 90000 and 100000
```

Tuple comparison

```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```

The from Clause

The **from** clause lists the relations involved in the query

Corresponds to the **Cartesian product** operation of the relational algebra.

Find the Cartesian product *instructor X teaches*

```
select *
from instructor, teaches
```

generates every possible instructor – teaches pair, with all attributes from both relations

Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra)

Joins

For all instructors who have taught some course, find their names and the course ID of the courses they taught.

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID
```

Find the course ID, semester, year and title of each course offered by the Comp. Sci. department

```
select section.course_id, semester, year, title
from section, course
where section.course_id = course.course_id and
dept_name = 'Comp. Sci.'
```



Natural Join Example

List the names of instructors along with the course ID of the courses that they taught.

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID;
```

```
select name, course_id
from instructor natural join teaches;
```

```
teaches(ID, course_id, sec_id, semester, year)
instructor(ID, name, dept_name, salary)
```

Natural Join

Natural join matches tuples with the same values for all common attributes, and retains only one copy of each common column

select * from instructor natural join teaches;

```
instructor(ID, name, dept_name, salary)
teaches(ID, course_id, sec_id, semester, year)
```

ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010

Natural Join (Cont.)

Beware of **unrelated attributes with same name** which get equated incorrectly

List the names of instructors along with the titles of courses that they teach

```
course(course_id, title, dept_name, credits)
teaches(ID, course_id, sec_id, semester, year)
instructor(ID, name, dept_name, salary)
```

Incorrect version (makes *course.dept_name = instructor.dept_name*)

```
select name, title
from instructor natural join teaches natural join course;
```

Correct version

```
select name, title
from instructor natural join teaches, course
where teaches.course_id = course.course_id;
select name, title
from (instructor natural join teaches) join course using(course_id);
select name, title
from instructor, teaches, course
where instructor.ID = teaches.ID and teaches.course_id = course.course_id;
```

Natural Join Another Example

Find students who takes courses across his/her department.

```
select distinct student.id
from (student natural join takes)
      join course using (course_id)
where student.dept_name <> course.dept_name
```

```
student( ID, name, dept_name, tot_cred )
takes ( ID, course_id, sec_id, semester, year, grade )
course( course_id, title, dept_name, credits )
```

The Rename Operation

The SQL allows renaming relations and attributes using the **as** clause:

old-name as new-name

E.g.

```
select ID, name, salary/12 as monthly_salary
from instructor
```

Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```

Keyword **as** is optional and may be omitted

instructor as T ≡ instructor T

Keyword **as** must be omitted in Oracle

String Operations

SQL includes a string-matching operator for comparisons on character strings. The operator "like" uses patterns that are described using two special characters:

- percent (%). The % character matches any substring.
- underscore (_). The _ character matches any character.

Find the names of all instructors whose name includes the substring "dar".

```
select name
from instructor
where name like '%dar%'
```

Match the string "100 %"

```
like '100 \%' escape '\'
like '100 \%''
like '100 #%' escape '#'
```

String Operations (Cont.)

Patters are case sensitive.

Pattern matching examples:

- 'Intro%' matches any string beginning with "Intro".
- '%Comp%' matches any string containing "Comp" as a substring.
- '___' matches any string of exactly three characters.
- '___ %' matches any string of at least three characters.

SQL supports a variety of string operations such as
concatenation (using "||")
converting from upper to lower case (and vice versa)
finding string length, extracting substrings, etc.

Ordering the Display of Tuples

List in alphabetic order the names of all instructors

```
select distinct name
from instructor
order by name
```

We may specify **desc** for descending(降) order or **asc** for ascending(升) order, for each attribute; ascending order is the default.

Example: **order by name desc**

Can sort on multiple attributes

Example: **order by dept_name, name**

The limit Clause

The **limit** clause can be used to constrain the number of rows returned by the select statement.

limit clause takes one or two numeric arguments, which must both be nonnegative integer constants:

```
limit offset, row_count
limit row_count == limit 0, row_count
```

List names of instructors whose salary is among top 3.

```
select name
from instructor
order by salary desc
limit 3; // limit 0,3
```

Duplicates

In relations with duplicates, SQL can define how many copies of tuples appear in the result.

Multiset (多重集) versions of some of the relational algebra operators – given multiset relations r_1 and r_2 :

1. $\sigma_\theta(r_1)$: If there are c_1 copies of tuple t_1 in r_1 , and t_1 satisfies selections σ_θ , then there are c_1 copies of t_1 in $\sigma_\theta(r_1)$.
2. $\Pi_A(r)$: For each copy of tuple t_1 in r_1 , there is a copy of tuple $\Pi_A(t_1)$ in $\Pi_A(r_1)$ where $\Pi_A(t_1)$ denotes the projection of the single tuple t_1 .
3. $r_1 \times r_2$: If there are c_1 copies of tuple t_1 in r_1 and c_2 copies of tuple t_2 in r_2 , there are $c_1 \times c_2$ copies of the tuple t_1, t_2 in $r_1 \times r_2$.

Duplicates (Cont.)

Example: Suppose multiset relations $r_1(A, B)$ and $r_2(C)$ are as follows:

$r_1 = \{(1, a) (2, a)\} \quad r_2 = \{(2), (3), (3)\}$

$\Pi_B(r_1) = \{(a), (a)\}$

$\Pi_B(r_1) \times r_2 = \{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$

SQL duplicate semantics:

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

is equivalent to the *multiset* version of the expression:

$\prod_{A_1, A_2, \dots, A_n} (\sigma_P(r_1 \times r_2 \times \dots \times r_m))$

Set Operations

Find courses that ran in Fall 2009 or in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)
union
(select course_id from section where sem = 'Spring' and year = 2010)
```

Find courses that ran in Fall 2009 and in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)
intersect
(select course_id from section where sem = 'Spring' and year = 2010)
```

Find courses that ran in Fall 2009 but not in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)
except
(select course_id from section where sem = 'Spring' and year = 2010)
```

Set Operations

Set operations **union**, **intersect**, and **except**

Each of the above operations automatically eliminates duplicates

To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

Suppose a tuple occurs m times in r and n times in s , then, it occurs:

$$\begin{aligned} & m + n \text{ times in } r \text{ union all } s \\ & \min(m, n) \text{ times in } r \text{ intersect all } s \\ & \max(0, m - n) \text{ times in } r \text{ except all } s \end{aligned}$$

Null Values

null signifies an unknown value or that a value does not exist.

The result of any arithmetic expression involving **null** is **null**

Example: $5 + \text{null}$ returns null

The predicate **is null** can be used to check for null values.

Example: Find all instructors whose salary is null.

```
select name
from instructor
where salary is null
```

Null Values

Comparisons with null values return the special truth value: **unknown**

Three-valued logic using the truth value **unknown**:

OR: $(\text{unknown or true}) = \text{true}$,
 $(\text{unknown or false}) = \text{unknown}$
 $(\text{unknown or unknown}) = \text{unknown}$

AND: $(\text{true and unknown}) = \text{unknown}$,
 $(\text{false and unknown}) = \text{false}$,
 $(\text{unknown and unknown}) = \text{unknown}$

NOT: $(\text{not unknown}) = \text{unknown}$

In SQL " P is **unknown**" evaluates to true if predicate P evaluates to **unknown**

Result of select predicate is treated as **false** if it evaluates to **unknown**

Aggregate Functions

These functions operate on the multiset of values of a column of a relation, and return a value

avg: average value
min: minimum value
max: maximum value
sum: sum of values
count: number of values

Aggregate Functions (Cont.)

Find the average salary of instructors in the Computer Science department

```
select avg (salary)
from instructor
where dept_name= 'Comp. Sci.';
```

Find the total number of instructors who teach a course in the Spring 2010 semester

```
select count (distinct ID)
from teaches
where semester = 'Spring' and year = 2010
```

Find the number of tuples in the *course* relation

```
select count (*)
from course;
```

Aggregate Functions – Group By

Find the average salary of instructors in each department

```
select dept_name, avg (salary)
from instructor
group by dept_name;
```

Note: departments with no instructor will not appear in result

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

Aggregation (Cont.)

Attributes in **select** clause outside of aggregate functions **must appear** in **group by** list

```
/* erroneous query */
select dept_name, ID, avg (salary)
from instructor
group by dept_name;
```

Aggregate Functions – Having Clause

Null Values and Aggregates

Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg(salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

```
select dept_name, count (*) as cnt
from instructor
where salary >=100000
group by dept_name
having count (*) > 10
order by cnt;
```

Total all salaries

```
select sum (salary)
from instructor
```

Above statement ignores null amounts

Result is **null** if there is no non-null amount

All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes

What if collection has only null values?

count returns 0

all other aggregates return null

Arithmetic expression with Aggregate functions

Nested Subqueries

Find departments in which there is no duplicate name of students.

```
select dept_name
from student
group by dept_name
having count(distinct name) = count(id)
```

What is the meaning of the following statement ?

```
select dept_name
from student
group by dept_name
having 1-count(distinct name)/ count(id)<0.001 ;
```

SQL provides a mechanism for the nesting of subqueries.

A **subquery** is a **select-from-where** expression that is nested within another query.

A common use of subqueries is to perform tests for :

set membership
set comparisons
set cardinality

Set Membership

Set Membership

Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

```
select count (distinct ID)
from takes
where (course_id, sec_id, semester, year) in
(select course_id, sec_id, semester, year
from teaches
where teaches.ID= '10101');
```

Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id not in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

Set Comparison

Set Comparison

Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Biology';
```

Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name
from instructor
where salary > all (select salary
from instructor
where dept_name = 'Biology');
```

Same query using > **some** clause

```
select name
from instructor
where salary > some (select salary
from instructor
where dept_name = 'Biology');
```

Scalar Subquery

Scalar (标量) subquery is one which is used where a single value is expected

E.g. `select name
from instructor
where salary * 10 >
(select budget from department
where department.dept_name = instructor.dept_name)`

Runtime error if subquery returns more than one result tuple

Test for Empty Relations

The **exists** construct returns the value **true** if the argument subquery is nonempty.

exists $r \Leftrightarrow r \neq \emptyset$

not exists $r \Leftrightarrow r = \emptyset$

[Skip to modification of database](#)

Correlation Variables

Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id  
from section as S  
where semester = 'Fall' and year = 2009 and  
exists (select *  
        from section as T  
        where semester = 'Spring' and year = 2010  
          and S.course_id = T.course_id);
```

Correlated subquery

Correlation name or correlation variable

Not Exists

Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name  
from student as S  
where not exists (select course_id  
                  from course  
                  where dept_name = 'Biology')  
except  
(select T.course_id  
  from takes as T  
  where S.ID = T.ID));
```

Note that $X - Y = \emptyset \Leftrightarrow X \subseteq Y$

Note: Cannot write this query using = all and its variants

Test for Absence of Duplicate Tuples

The **unique** construct tests whether a subquery has any duplicate tuples in its result.

(Evaluates to “true” on an empty set)

Find all courses that were offered at most once in 2009

```
select T.course_id  
from course as T  
where unique (select R.course_id  
              from section as R  
              where T.course_id = R.course_id  
                and R.year = 2009);
```

Test for Absence of Duplicate Tuples

Find all courses that were offered once in 2009

```
select T.course_id  
from course as T  
where unique (select R.course_id  
              from section as R  
              where T.course_id = R.course_id  
                and R.year = 2009)  
and exists (select R.course_id  
              from section as R  
              where T.course_id = R.course_id  
                and R.year = 2009);
```

Another solution:

```
and course_id in (select course_id  
                   from section  
                   where year = 2009);
```

Test for Absence of Duplicate Tuples

Find all courses that were offered at most once in every semester

```
select T.course_id  
from course as T  
where unique (select R.course_id, year, semester  
              from section as R  
              where T.course_id = R.course_id  
                );
```

Find all courses that were offered once in every semester

```
select T.course_id  
from course as T  
where unique (select R.course_id, year, semester  
              from section as R  
              where T.course_id = R.course_id  
                )  
and exists (select R.course_id, year, semester  
              from section as R  
              where T.course_id = R.course_id  
                );
```

*With Clause

The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the **with** clause occurs.

Find all departments with the maximum budget

```
with max_budget (value) as  
(select max(budget)  
  from department)  
select dept_name  
from department, max_budget  
where department.budget = max_budget.value;
```

```
select dept_name  
from department  
where budget = (select (max(budget)) from department))
```

*Complex Queries using With Clause

With clause is very useful for writing complex queries

Supported by most database systems, with minor syntax variations

Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total(dept_name, value) as
  (select dept_name, sum(salary)
   from instructor
   group by dept_name),
  dept_total_avg(value) as
  (select avg(value)
   from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value >= dept_total_avg.value;
```

Modification of the Database

Deletion of tuples from a given relation

Insertion of new tuples into a given relation

Updating values in some tuples in a given relation

Modification of the Database – Deletion

Delete all instructors

```
delete from instructor
```

Delete all instructors from the Finance department

```
delete from instructor
where dept_name = 'Finance';
```

Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

```
delete from instructor
where dept_name in (select dept_name
                     from department
                     where building = 'Watson');
```

Deletion (Cont.)

Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor
where salary < (select avg(salary) from instructor);
```

Problem: as we delete tuples from deposit, the average salary changes

Solution used in SQL:

1. First, compute **avg** salary and find all tuples to delete
2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

Modification of the Database – Insertion

Add a new tuple to *course*

```
insert into course
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

or equivalently

```
insert into course(course_id, title, dept_name, credits)
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

Add a new tuple to *student* with *tot_creds* set to null

```
insert into student
values ('3003', 'Green', 'Finance', null);
```

Insertion (Cont.)

Add all instructors to the *student* relation with *tot_creds* set to 0

```
insert into student
select ID, name, dept_name, 0
from instructor
```

The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

otherwise queries like

```
insert into table1 select * from table1
```

would cause problems, if *table1* did not have any primary key defined.

Modification of the Database – Updates

Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others receive a 5% raise

Write two **update** statements:

```
update instructor
set salary = salary * 1.03
where salary > 100000;
update instructor
set salary = salary * 1.05
where salary <= 100000;
```

The order is important

Can be done better using the **case** statement (next slide)

Case Statement for Conditional Updates

Same query as before but with case statement

```
update instructor
set salary = case
  when salary <= 100000 then salary * 1.05
  else salary * 1.03
end
```

Updates with Scalar Subqueries

Recompute and update tot_creds value for all students

```
update student S
  set tot_cred = ( select sum(credits)
                    from takes natural join course
                   where S.ID= takes.ID and
                         takes.grade <> 'F' and
                         takes.grade is not null);
```

Sets tot_creds to null for students who have not taken any course

Instead of `sum(credits)`, use:

```
case
  when sum(credits) is not null then sum(credits)
  else 0
end
```

End of Chapter 3