

Joined Relation
SQL Data Types and Schemas
Integrity Constraints
Views
Indexes
Transactions
Authorization

Chapter 4: Intermediate SQL

Joined Relations

Join operations take two relations and return as a result another relation.

Join operations are typically used as subquery expressions in the **from** clause

Join condition – defines which tuples in the two relations match, and what attributes are present in the result of the join.

Join type – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

Join types	Join Conditions	
inner join	natural	← 自然连接
left outer join	on <predicate>	← 条件连接
right outer join	using (A_1, A_1, \dots, A_n)	← 等值连接
full outer join		

Outer Join

An extension of the join operation that avoids loss of information.

Computes the join and then adds tuples form one relation that does not match tuples in the other relation to the result of the join.

Uses *null* values.

Right Outer Join

course **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

```
select count(*)
from course natural right outer join prereq
where prereq_id is null;
```

Join operations – Example

Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Observe that
prereq information is missing for CS-315 and
course information is missing for CS-437

Left Outer Join

course **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

Full Outer Join

course **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

Joined Relations – Examples

course inner join prereq on
course.course_id = prereq.course_id

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

What is the difference between the above, and a natural join?

course left outer join prereq on
course.course_id = prereq.course_id

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null

Built-in Data Types in SQL

date: Dates, containing a (4 digit) year, month and date

Example: **date** ‘2005-7-27’

time: Time of day, in hours, minutes and seconds.

Example: **time** ‘09:00:30’ **time** ‘09:00:30.75’

timestamp: date plus time of day

Example: **timestamp** ‘2005-7-27 09:00:30.75’

interval: period of time

Example: interval ‘1’ day

Subtracting a date/time/timestamp value from another gives an interval value

Interval values can be added to date/time/timestamp values

date, time functions:

current_date(), current_time()

year(x), month(x), day(x), hour(x), minute(x), second(x)

Domains

create domain construct in SQL-92 creates user-defined domain types

create domain person_name char(20) not null

Types and domains are similar. Domains can have constraints, such as not null, specified on them.

create domain degree_level varchar(10)
constraint degree_level_test
check (value in ('Bachelors', 'Masters', 'Doctorate'));

Integrity Constraints

Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.

A checking account must have a balance greater than \$10,000.00

A salary of a bank employee must be at least \$4.00 an hour

A customer must have a (non-null) phone number

Joined Relations – Examples

course natural right outer join prereq

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

course full outer join prereq using (course_id)

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101

User-Defined Types

create type construct in SQL creates user-defined type

create type Dollars as numeric (12,2) final

create table department
(dept_name varchar (20),
building varchar (15),
budget Dollars);

Large-Object Types

Large objects (photos, videos, CAD files, etc.) are stored as a large object.

blob: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)

MySQL BLOB datatypes:

- ▶ **TinyBlob** : 0 ~ 255 bytes.
- ▶ **Blob** : 0 ~ 64K bytes.
- ▶ **MediumBlob** : 0 ~ 16M bytes.
- ▶ **LargeBlob** : 0 ~ 4G bytes.

clob: character large object -- object is a large collection of character data

When a query returns a large object, a pointer is returned rather than the large object itself.

Integrity Constraints on a Single Relation

not null

primary key

unique

check (P), where P is a predicate

foreign key

Not Null and Unique Constraints

not null

Declare *name* and *budget* to be **not null**
name **varchar**(20) **not null**
budget **numeric**(12,2) **not null**

unique (A₁, A₂, ..., A_m)

The unique specification states that the attributes A1, A2, ... Am form a super key (✕ **candidate key**) .
Candidate keys are permitted to be null (in contrast to primary keys).

Referential Integrity

Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.

Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology”.

Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S.

Integrity Constraint Violation During Transactions

```
create table person (  
  ID char(10),  
  name char(40),  
  mother char(10),  
  father char(10),  
  primary key (ID),  
  foreign key (father) references person,  
  foreign key (mother) references person);
```

How to insert a tuple without causing constraint violation ?
insert father and mother of a person before inserting person
OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)
OR defer constraint checking to **transaction end**.

Complex Check Clauses

Unfortunately: subquery in check clause not supported by pretty much any database
Alternative: **triggers**

The check clause

check (P)
where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (  
  course_id varchar (8),  
  sec_id varchar (8),  
  semester varchar (6),  
  year numeric (4,0),  
  building varchar (15),  
  room_number varchar (7),  
  time slot id varchar (4),  
  primary key (course_id, sec_id, semester, year),  
  check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```

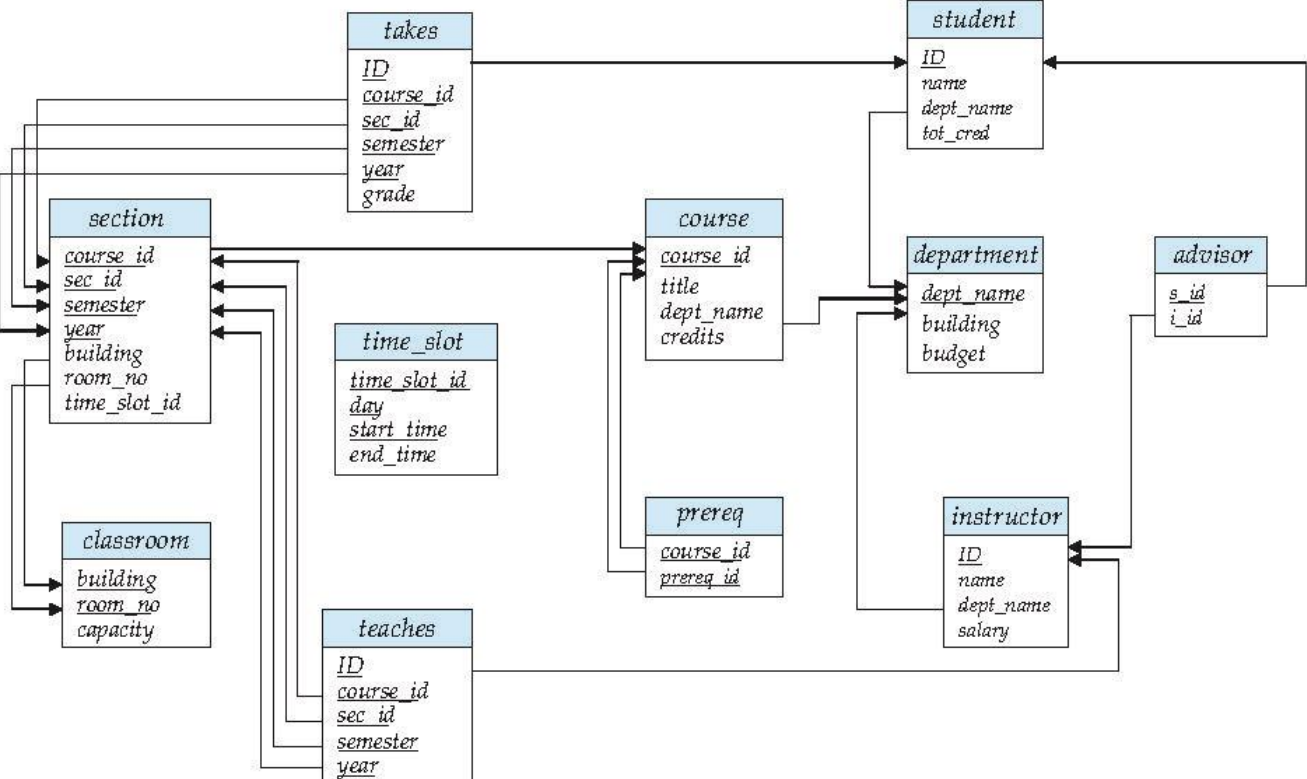
Cascading Actions in Referential Integrity

```
create table course (  
  course_id char(5) primary key,  
  title varchar(20),  
  dept_name varchar(20) references department  
)  
create table course (  
  ...  
  dept_name varchar(20),  
  foreign key (dept_name) references department  
    on delete cascade  
    on update cascade,  
  ...  
)
```

alternative actions to cascade: **set null, set default, restricted**

Complex Check Clauses

```
check (time_slot_id in  
  (select time_slot_id from time_slot))
```



Every section has at least one instructor teaching the section.

```
check ((course_id, sec_id, semester, year) in  
  (select course_id, sec_id, semester, year from teaches))
```

assertion

```
create assertion <assertion-name> check  
<predicate>;  
create assertion credits_earned_constraint check  
(not exists  
  (select ID  
    from student  
    where tot_cred <> (  
      select sum(credits)  
        from takes natural join course  
        where student.ID=takes.ID  
          and grade is not null  
          and grade<>'F'))
```


Views

A **view** provides a mechanism to hide certain data from the view of certain users.

Consider a person who needs to know [an instructors name and department, but not the salary](#). This person should see a relation described, in SQL, by

```
select ID, name, dept_name
from instructor
```

Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.

Example Views

A view of instructors without their salary

```
create view faculty as
select ID, name, dept_name
from instructor
```

Find the names of all instructors in the Biology department

```
select name
from faculty
where dept_name = 'Biology'
```

Create a view of department salary totals

```
create view departments_total_salary(dept_name, total_salary) as
select dept_name, sum (salary)
from instructor
group by dept_name;
```

View Expansion

Expand use of a view in another **view**

```
create view physics_fall_2009_watson as
select course_id, room_number
from physics_fall_2009
where building= 'Watson';
→
create view physics_fall_2009_watson as
(select course_id, room_number
from (select course.course_id, building, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009')
where building= 'Watson');
```

Update of a View

Add a new tuple to *faculty* view which we defined earlier

```
create view faculty as
select ID, name, dept_name
from instructor
```

insert into faculty values ('30765', 'Green', 'Music');

This insertion must be represented by the insertion of the tuple
('30765', 'Green', 'Music', null)

into the *instructor* relation

insert into instructor values ('30765', 'Green', 'Music', null);

View Definition

A view is defined using the **create view** statement which has the form

```
create view v as < query expression >
```

where <query expression> is any legal SQL expression. The view name is represented by *v*.

Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.

View definition is not the same as creating a new relation by evaluating the query expression

Rather, **a view definition causes the saving of an expression; the expression is substituted into queries using the view.**

Views Defined Using Other Views

```
create view physics_fall_2009 as
select course.course_id, sec_id, building, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009';
```

```
create view physics_fall_2009_watson as
select course_id, room_number
from physics_fall_2009
where building= 'Watson';
```

View Expansion

Expand use of a view in a **query**:

```
select course_id, room_number
from physics_fall_2009
where building= 'Watson';
→
select course_id, room_number
from (select course.course_id, building, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009')
where building= 'Watson';
→
select course_id, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009'
and building= 'Watson';
```

Some Updates cannot be Translated Uniquely

```
create view instructor_info as
select ID, name, building
from instructor, department
where instructor.dept_name= department.dept_name;
insert into instructor_info values ('69987', 'White', 'Taylor');
```

► which department, if multiple departments in Taylor?

Most SQL implementations allow updates only on simple views(**updatable views**)

The **from** clause has only one database relation.

The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.

Any attribute not listed in the **select** clause can be set to null

The query does not have a **group** by or **having** clause.

*Materialized Views

Materializing a view: create a physical table containing all the tuples in the result of the query defining the view

If relations used in the query are updated, the materialized view result becomes out of date

Need to maintain the view, by updating the view whenever the underlying relations are updated.

Example:

```
create materialized view departments_total_salary(dept_name, total_salary) as
select dept_name, sum (salary)
from instructor
group by dept_name;

select dept_name
from departments_total_salary
where total_salary > (select avg(total_salary) from departments_total_salary );
```

Indexes

```
create table student
(
    ID varchar (5),
    name varchar (20) not null,
    dept_name varchar (20),
    tot_cred numeric (3,0) default 0,
    primary key (ID) )
```

```
create index studentID_index on student(ID)
```

Indices are data structures used to speed up access to records with specified values for index attributes

```
e.g. select *
from student
where ID = '12345'
```

can be executed by using the index to find the required record, without looking at all records of student

* View and Logical Data Indepence

If relation S(a, b, c) is split into two sub relations S1(a,b) and S2(a,c)
How to realize the logical data independence?

- 1) create table S1 ...;
create table S2 ...;
- 2) insert into S1 select a, b from S;
insert into S2 select a, c from S;
- 3) drop table S;
- 4) create view S(a,b,c) as select a,b,c from S1 natural join S2;

select * from S where ... → select * from S1 natural join S2 where ...

insert into S values (1 ,2,3) → insert into S1 values (1, 2);
insert into S2 values (1 ,3);

Transactions

- Unit of work (NONE or ALL)
- Atomic transaction
 - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
 - Ended by commit work or rollback work
- But default on most databases: each SQL statement commits automatically
 - Can turn off auto commit for a session (e.g. using API)
- In MySQL:
 - >SET AUTOCOMMIT=0;
- In SQL:1999, can use: begin atomic end
 - Not supported on most databases

Transactions

```
Transaction example :
SET AUTOCOMMIT=0;
UPDATE account SET balance=balance -100 WHERE ano='1001';
UPDATE account SET balance=balance+100 WHERE ano='1002';
COMMIT;

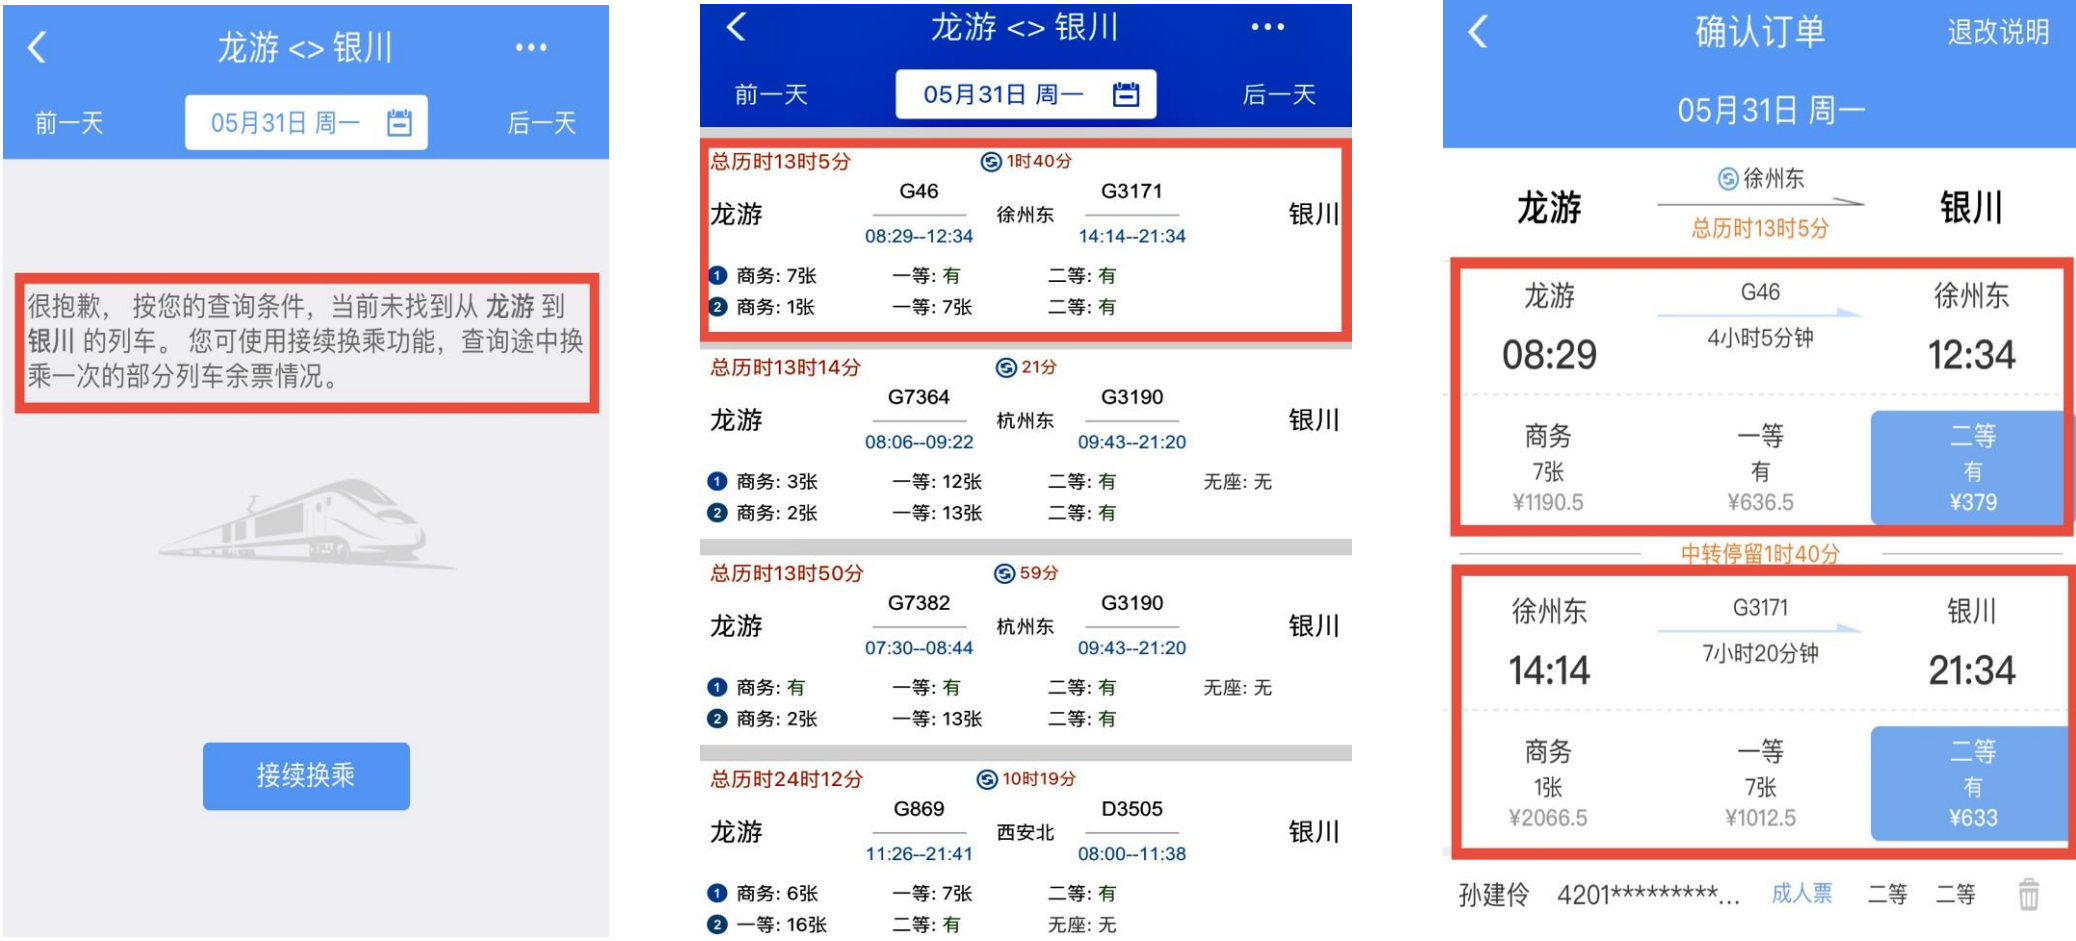
UPDATE account SET balance=balance -200 WHERE ano='1003';
UPDATE account SET balance=balance+200 WHERE ano='1004';
COMMIT;

UPDATE account SET balance=balance+balance*2.5%;
COMMIT;
```

Transaction Examples



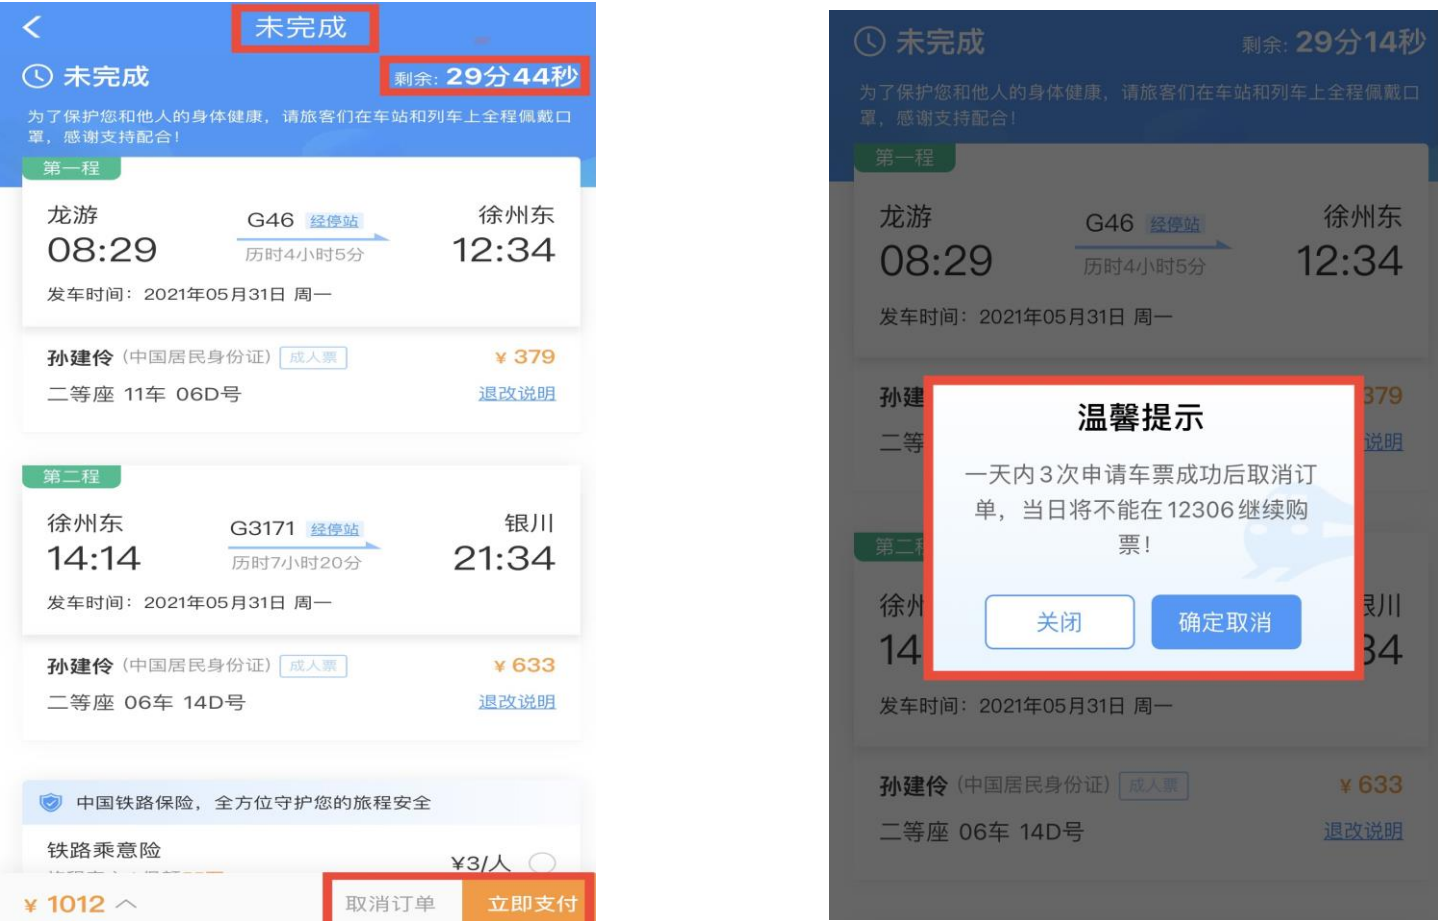
Transaction Boundaries
Booking one ticket vs multiple tickets?



Transaction Examples



Transaction Boundaries
Booking tickets & paying order - combining vs separating ?



ACID Properties

A transaction is a unit of program execution that accesses and possibly updates various data items.To preserve the integrity of data the database system must ensure:

- Atomicity. Either all operations of the transaction are properly reflected in the database or none are.
- Consistency. Execution of a transaction in isolation preserves the consistency of the database.
- Isolation. Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
- Durability. After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

Authorization(授权)

Forms of authorization on parts of the database:

- Select** - allows reading, but not modification of data.
- Insert** - allows insertion of new data, but not modification of existing data.
- Update** - allows modification, but not deletion of data.
- Delete** - allows deletion of data.

Forms of authorization to modify the database schema

- Resources (MySQL: Create)** - allows creation of new relations.
- Alteration** - allows addition or deletion of attributes in a relation.
- Drop** - allows deletion of relations.
- Index** - allows creation and deletion of indices.
- Create view (MySQL)** – allows creation of views.

Privileges in SQL

```
grant select on instructor to U1, U2, U3
grant select on department to public
grant update (budget) on department to U1,U2
grant all privileges on department to U1
```

Roles

```
create role instructor;
grant instructor to Amit;
```

Privileges can be granted to roles:


```
grant select on takes to instructor;
```

Roles can be granted to users, as well as to other roles

```
create role teaching_assistant
grant teaching_assistant to instructor;
  ▶ Instructor inherits all privileges of teaching_assistant
```

Chain of roles

```
create role dean;
grant instructor to dean;
grant dean to Satoshi;
```



Other Authorization Features

references privilege to create foreign key

```
grant reference (dept_name) on department to Mariano;
```

why is this required?



Authorization Specification in SQL

The **grant** statement is used to confer authorization

```
grant <privilege list> // privilege: 权限
on <relation name or view name> to <user list>
```

<user list> is:

- a user-id
- public**, which allows all valid users the privilege granted
- A role (more on this later)

Granting a privilege on a view does not imply granting any privileges on the underlying relations.

The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

Revoking Authorization in SQL

The **revoke** statement is used to revoke authorization.

```
revoke <privilege list>
on <relation name or view name>
from <user list>
```

Example:

```
revoke select on branch from U1, U2, U3
```

<privilege-list> may be **all** to revoke all privileges the revokee may hold.

If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.

If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.

All privileges that depend on the privilege being revoked are also revoked.


Authorization on Views

```
create view geo_instructor as
(select *
from instructor
where dept_name = 'Geology');
grant select on geo_instructor to geo_staff
```

Other Authorization Features

transfer of privileges

```
grant select on department to Amit with grant option;
revoke select on department from Amit, Satoshi cascade;
revoke select on department from Amit, Satoshi restrict;
revoke grant option for select on department from Amit;
```



End of Chapter 4